

You will see the features {} with sequences of 50 text generations each. Each sequence consists of an original text and a modified version where a specific hidden feature has been gradually strengthened or weakened. The same hidden feature is shifted consistently across all sequences.

Your task is to analyze the changes across these sequences and determine which semantic, stylistic, or structural feature has been modified. Try to find for each feature the dependencies and hidden meaning.

Output Format:

Create a structured table with the following columns:

Feature Number: A unique identifier for the observed feature.

Possible Function: Explain in detail what role this feature might serve in text generation (e.g., enhancing coherence, increasing formality, affecting emotional tone).

Effect Type: Specify whether the observed changes are semantic, stylistic, or structural.

Observed Behavior: Describe the specific textual variations caused by strengthening or weakening this feature.

Each row should correspond to a distinct feature, listing its effects and possible functions with sufficient explanation

Figure 17: Prompt used for steering analysis