

Limitation and broader impact. A main limitation of this work is that detector-evasion results are unlikely to remain static. Commercial detectors can adapt in response to known attacks, so a method that works today may weaken as detectors are updated. The broader impact is mixed. On the one hand, HIP could be misused to evade deployed detectors. On the other hand, that risk already exists in practice, and the paper makes it more explicit by identifying a concrete failure mode on current commercial systems. We view the main contribution as diagnostic: identifying this weakness should help detector researchers build stronger methods against paraphrase-based evasion.

Acknowledgments

This work was supported in part by NSF grant IIS-2046640 (CAREER).

This work was also supported by GPTZero and Pangram’s research credits.

References

- George Alexandru Adam, Alexander Cui, Edwin Thomas, Emily Napier, Nazar Shmatko, Jacob Schnell, Jacob Junqi Tian, Alekhya Dronavalli, Edward Tian, and Dongwon Lee. Gptzero: Robust detection of llm-generated texts. *arXiv preprint arXiv:2602.13042*, 2026.
- Bradley Emi and Max Spero. Technical report on the pangram ai-generated text classifier. *arXiv preprint arXiv:2402.14873*, 2024.
- GPTZero. Ai detector for teachers & educators. <https://gptzero.me/educators>, 2026. Accessed: 2026-04-27.
- Pangram. Ai detector for teachers & educators. <https://www.pangram.com/use-cases/teachers>, 2026. Accessed: 2026-04-27.
- Isaac David and Arthur Gervais. Authormist: Evading ai text detectors with reinforcement learning. *arXiv preprint arXiv:2503.08716*, 2025.
- Suraj Ranganath and Atharv Ramesh. Stealthrl: Reinforcement learning paraphrase attacks for multi-detector evasion of ai-text detectors. *arXiv preprint arXiv:2602.08934*, 2026.
- Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John Wieting, and Mohit Iyyer. Paraphrasing evades detectors of ai-generated text, but retrieval is an effective defense. *Advances in neural information processing systems*, 36:27469–27500, 2023.
- Aldan Creo and Shushanta Pudasaini. Silverspeak: evading ai-generated text detectors using homographs. In *Proceedings of the 1st Workshop on GenAI Content Detection (GenAIDetect)*, pages 1–46, 2025.
- Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D Manning, and Chelsea Finn. Detectgpt: Zero-shot machine-generated text detection using probability curvature. In *International conference on machine learning*, pages 24950–24962. PMLR, 2023.
- Abhimanyu Hans, Avi Schwarzschild, Valeriia Cherepanova, Hamid Kazemi, Aniruddha Saha, Micah Goldblum, Jonas Geiping, and Tom Goldstein. Spotting llms with binoculars: Zero-shot detection of machine-generated text. *arXiv preprint arXiv:2401.12070*, 2024.
- Katherine Thai, Bradley Emi, Elyas Masrour, and Mohit Iyyer. Editlens: Quantifying the extent of ai editing in text. *arXiv preprint arXiv:2510.03154*, 2025.
- Qihui Zhang, Chujie Gao, Dongping Chen, Yue Huang, Yixin Huang, Zhenyang Sun, Shilin Zhang, Weiye Li, Zhengyan Fu, Yao Wan, et al. Llm-as-a-coauthor: Can mixed human-written and machine-generated text be detected? In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 409–436, 2024.
- Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémy Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, et al. Open problems and fundamental limitations of reinforcement learning from human feedback. *arXiv preprint arXiv:2307.15217*, 2023.
- Prasann Singhal, Tanya Goyal, Jiacheng Xu, and Greg Durrett. A long way to go: Investigating length correlations in rlhf. In *First Conference on Language Modeling*, 2024.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R Bowman, Esin Durmus, Zac Hatfield-Dodds, Scott R Johnston, Shauna M Kravec, et al. Towards understanding sycophancy in language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- Rajiv Movva, Smitha Milli, Sewon Min, and Emma Pierson. What’s in my human feedback? learning interpretable descriptions of preference data, 2026. URL <https://arxiv.org/abs/2510.26202>.
- Ishan Jindal, Pranjal Bharti, Lakkidi Vinay, SACHIN DEV SHARMA, et al. Keep the alignment, skip the overhead: Lightweight instruction alignment for continually trained llms. In *Second Workshop on Test-Time Adaptation: Putting Updates to the Test! at ICML 2025*, 2025.

- Jack Morris. Introducing gpt-oss-20b-base. Post on X (formerly Twitter), August 2025. URL <https://x.com/jxmnp/status/1955436067353502083>. Accessed: 2026-05-06.
- Vinu Sankar Sadasivan, Aounon Kumar, Sriram Balasubramanian, Wenxiao Wang, and Soheil Feizi. Can ai-generated text be reliably detected? *arXiv preprint arXiv:2303.11156*, 2023.
- Elyas Masrour, Bradley N Emi, and Max Spéro. Damage: detecting adversarially modified ai generated text. In *Proceedings of the 1st Workshop on GenAI Content Detection (GenAIDetect)*, pages 120–133, 2025.
- Junjie Huang, Ruiquan Zhang, Jinsong Su, and Yidong Chen. Tempparaphraser: “heating up” text to evade ai-text detection through paraphrasing. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 31542–31561, 2025.
- Yongtong Gu, Songze Li, and Xia Hu. Mash: Evading black-box ai-generated text detectors via style humanization. *arXiv preprint arXiv:2601.08564*, 2026.
- Jingyi Zheng, Junfeng Wang, Zhen Sun, Wenhan Dong, Yule Liu, and Xinlei He. Th-bench: Evaluating evading attacks via humanizing ai text on machine-generated text detectors. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, pages 5948–5959, 2025.
- Yiwei Zha, Rui Min, and Shanu Sushmita. Padben: A comprehensive benchmark for evaluating ai text detectors against paraphrase attacks. *arXiv preprint arXiv:2511.00416*, 2025.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. *Advances in neural information processing systems*, 36:46534–46594, 2023.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Liang Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *Iclr*, 1(2):3, 2022.
- Liam Dugan, Alyssa Hwang, Filip Trhlík, Andrew Zhu, Josh Magnus Ludan, Hainiu Xu, Daphne Ippolito, and Chris Callison-Burch. Raid: A shared benchmark for robust evaluation of machine-generated text detectors. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12463–12492, 2024.
- Yafu Li, Qintong Li, Leyang Cui, Wei Bi, Zhilin Wang, Longyue Wang, Linyi Yang, Shuming Shi, and Yue Zhang. Mage: Machine-generated text detection in the wild. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 36–53, 2024.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. *Advances in neural information processing systems*, 36:10088–10115, 2023.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th symposium on operating systems principles*, pages 611–626, 2023.
- Guilherme Penedo, Hynek Kydlíček, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro Von Werra, Thomas Wolf, et al. The fineweb datasets: Decanting the web for the finest text data at scale. *Advances in Neural Information Processing Systems*, 37:30811–30849, 2024.